

September 23, 2026

To the Hiring Committee
Anthropic

Re: Letter of recommendation for Atisha Rajpurohit

Dear Hiring Committee,

One Sunday evening this June, at the tail end of a nine-hour bootcamp day, I handed the floor to a teaching assistant and told a room full of engineers, “You’re looking at my successor.” The room laughed, as it was meant to. I was half joking. This letter is my attempt to explain why I no longer am.

A word about who is writing, so you can calibrate what follows. I run SupportVectors, an applied AI research and training lab and school in Fremont, where we teach working engineers everything from the foundations of AI through LLMs, fine-tuning, RAG, and agentic systems. Behind me are more than two decades as an enterprise and chief architect building large-scale ML platforms, fifteen years as a VP and SVP of Engineering, and, before that, multiple post-graduate degrees in computer science and theoretical particle physics. I have hired, mentored, and worked alongside many strong engineers. I have watched very few of them learn to teach.

I should also say what kind of place SupportVectors is, because it matters to what follows. It is not a place for casual training. Every bootcamp spans the full range, from reading research papers in depth, equation by equation, to building hands-on, enterprise-grade capstone projects, and we are considered one of the more rigorous places to learn AI. Our students are working engineers, and they are not easily impressed. In such a place, for someone so early in her career to rise into a teacher that room trusts is remarkable. Atisha has done it, and I have had the privilege of watching it happen up close, week by week—and, because we record and transcribe every session, on the record.

She came to us after completing her MS in Data Science at Georgetown University, and then did something unusual: she took our courses one after another, all of them. She was not collecting certificates. I first noticed her as the student who does the reading and then holds the instructor to it. When I opened one session by asking the class to recall the previous week’s six-dimensional taxonomy of agents, it was her summary, complete and exact, that made me stop and praise it in front of everyone. Later that same day she caught me in a contradiction: “You just said it’s not domain adaptation—but that’s how you started.” She was right. I had blurred two different senses of fine-tuning, and her question forced me to draw, on the spot, a distinction the whole room needed. Good students absorb; the rare ones push back, precisely and without fear. That is the soil a teacher grows from.

Across the following cohorts, Atisha moved naturally from repeat student to volunteer teaching assistant, and then into what I can only describe as a co-instructor. This summer’s AI Agents bootcamp made that unmistakable. On prompt-engineering day, she led the in-class experiments—the games where real learning happens in my classroom, and which are unforgiving to facilitate. On our hands-on Skills day, when much of the room ran into problems with shells, zip files, and Windows quirks, she sat beside the students who were stuck and quietly helped them move forward, ensuring that no one left feeling left behind. Over the twelve weeks, she also delivered six research talks of her own: on context engineering and context rot; MCP security, including a sixteen-threat taxonomy of tool misuse; the ACE paper; Memo; Anthropic’s harness-engineering articles; and, on the final day, OpenClaw and the placement of the trust boundary in personal agents. Most of these talks took place in the evening slot—7 PM,

sometimes closer to 8, after nine hours of lecture—the most expensive attention of the week and the hardest room any teacher will face. The room stayed. At the end of the Memo talk, she admitted she had feared her audience would run away. They had not. She had earned their attention.

Let me be specific about the craft, because “good speaker” is not what I am recommending.

She teaches the problem until the solution feels inevitable. In her ACE talk she spent her best minutes on the failure modes—brevity bias, and context collapse, where a carefully accumulated playbook of roughly 18,000 tokens collapses to about 120 in a single rewrite—so that by the time the architecture arrived, the room was already asking for it. To contrast ACE’s incremental updates with GEPA’s rewrite-from-memory she invented her own image, a Post-it note added to the page rather than the page being rewritten, and it worked better than anything I would have reached for. It was hers.

She has the Socratic instinct, and nobody drilled it into her. She opened the Memo talk with a game: here are the items from an air-crash investigation—where does each one go, memory, RAG, skills, or the context window? When the room answered “RAG” for the aircraft’s maintenance history, she did not correct them. She offered a counterexample—a doctor and a patient’s medical history—and asked, “Have I got you all to change your answer?” That is exactly how I teach, and it is the hardest thing to teach a teacher.

She thinks like an engineer, with critical distance. She presented the Anthropic harness articles as a fossil record of failures: every component introduced as the failure it was built to answer—context anxiety, premature victory, lenient self-grading—and she built an interactive artifact in which you click any component and see its failure and its fix. She was careful to scope every claim (“it worked at that time, for that goal, with that model”), a habit many senior speakers never acquire. And when a student asked her to separate context from harness, she answered, “The context is the object it works on; the harness is the doer.” I have read published definitions that are worse.

She goes and runs the experiment. For her OpenClaw talk she mapped the system onto the four tiers of memory the class already owned, laid out the trust-boundary spectrum from maximal to minimal as the real design question, and then ran a live experiment with a browser-use agent on her own account and her own money, reporting back failure rates, costs, and the verdict—accuracy yes, quality and latency not yet. Her root-cause diagnosis of the 2026 security incidents—“there was a harness that proposed and disposed, but there was almost no disposal policy”—is the best single line of that day, and my notes for the forthcoming textbook chapter now read: credit Atisha by name. For her MCP-security talk she read the paper just two days beforehand and still arrived with a hand-built interactive presentation the class talked about for weeks.

She is calibrated about the edge of what she knows. Asked something she was unsure of, she has never once bluffed. “I will look into it and put it on Discord. Promise.”—and then she does. In a field that runs on confident-sounding claims, a teacher who models where her knowledge ends is worth more than one who never seems to reach it. I suspect Anthropic, of all places, understands why.

And she learns faster than anyone I have coached. After each talk I gave her written feedback, sometimes blunt: referee the paper’s terminology against ours before you start; own every number on your slides; walk one path through the big diagram. Her next talk addressed most of it, unprompted. I have mentored principal architects from very large companies who took a year to absorb what she absorbed in a week. She is still sharpening—citing figures with exact provenance, holding the pause after her own question a beat longer—and she knows it, because I told her, and she is closing those gaps faster than I can list them. Her actuarial training shows up as a constant instinct to measure; nobody on my teaching staff is better placed to take apart an evaluation benchmark quantitatively.

Her reach goes well beyond my classroom. She has given public paper-reading talks at our Meetup; she organizes technical events and panel discussions around Silicon Valley, and she is an organizing member of the AI Collective. In other words, she can build a community, not just hold a room—and a training organization needs both.

I have long believed that teaching is a bridge between two loves: a love of the subject, and a love of the people with whom you share it. Many excellent engineers have the first. Atisha has both, and you can see the second in how she reads a tired room, sits beside the student who is stuck, and thanks people for staying. The role you describe asks for an AI engineer with a flair for teaching. I would put it the other way around: she is a teacher who also happens to be a first-rate AI engineer, and that ordering is exactly what makes her so rare.

The tradition I grew up in holds that a teacher's real reward is a student who carries the lamp further than he could. In June I told a room she was my successor, and they laughed. If you hire her, I will lose her—and I will be proud of the loss. I recommend her to you without reservation, and in the strongest terms I know. I would be happy to speak with anyone on the committee at any time.

Warm regards,

Asif Qamar

Founder, SupportVectors

asif@supportvectors.com · 510-298-9030