

Atisha Rajpurohit

 [Linkedin](#)  [Github](#)  [Medium](#)  atisha.rajpurohit@gmail.com  +1 (571)-351-3630

AI Engineer and statistician who builds production AI systems — and converts them, on the go, into things people can learn from: classes, podcasts, articles, and interactive teaching artifacts. First hands-on generative-AI instructor for a 60+ person consulting team; weekly lecturer to a 50–60 person cohort of senior leaders, TPMs, and engineers; designed and directed a five-event program reaching 400+ attendees. Runs an AI-assisted content pipeline (Claude, Claude Code, Claude Design) that turns research papers into teaching artifacts and working prototypes every week. I understand what needs to be understood, and how to teach it to be understood.

PROFESSIONAL EXPERIENCE

Cloudely (San Jose, CA) — *AI Engineer* Dec 2025 – Present

- Eliminated Konfeeg’s specialist bottleneck — schema and query changes required an admin — by building an AI application layer (FastAPI, Next.js, CI/CD) that lets non-technical users make app updates in natural language; led a 2–3 engineer team, improving development velocity 40%.
- Removed per-tenant consultant dependency — each customer’s data model needed bespoke, non-transferable training — by designing a LangGraph multi-agent system that plans against each tenant’s live schema with confirm-before-write gating, turning one-off consulting into a repeatable self-serve workflow.
- Reduced confidently-wrong answers on customers’ own data — where wrong answers erode enterprise trust — by engineering a hybrid retrieval architecture (BM25 + dense embeddings + Reciprocal Rank Fusion), measuring retrieval precision and contextual grounding up 35% and cutting clarification round-trips.
- Hosted and moderated the Cloudely Conversations podcast (“[From Data Chaos to Confident Decisions: Humans, Strategy & AI](#),” “[The AI Precipice: The Pulse Beneath](#)”), translating enterprise AI adoption and decision-making for non-technical business audiences.

Ernst & Young (Washington, D.C.) — *Quantitative Analyst* Jun 2022 – Oct 2025

- Built AI fluency across a 60+ person consulting team with no engineering background — Copilot, prompt engineering, agentic workflows — as their first hands-on instructor.
- Simulated statistical samples to validate model outputs and close historical data gaps across 125+ engagements, cutting predictive variability ~3% to meet strict compliance thresholds.
- Presented technical work and processes directly to senior leadership, building stakeholder trust and establishing clear ownership for a lean team.

Mercer (Mumbai) — *Data Analysis Manager*, global HR and benefits consultancy Mar 2018 – Mar 2021

- Owned client scoping and workflow planning for a 150+ account portfolio inside a people-function business — resolving ambiguous requests into actionable delivery plans — sustaining a 95% SLA target.
- Built predictive models to explain liability-movement drivers across 110+ client implementations, cutting unexplained liabilities 93% to meet strict regulatory thresholds.

TECHNICAL ENABLEMENT & TEACHING

Assistant Lecturer, *SupportVectors AI Labs* Jun 2026 – Present

- Teach weekly sessions on agent-engineering research to a 50–60 person cohort of senior leaders, technical program managers, and engineers — an arc spanning prompting and evaluation (promptfoo), MCP security, agentic context engineering (ACE), agent memory (Mem0), and harness engineering — using live demos, code prototypes, and interactive architecture diagrams.
- Build self-service teaching artifacts: an interactive MCP threat matrix the audience explores hands-on, architecture diagrams tracing a query’s journey through Mem0, and research papers reworked into decks and live-coded prototypes.
- Run an AI-assisted content pipeline — Claude, Claude Code, Claude Design — turning research papers into class-ready artifacts and prototypes on a weekly cadence.

Silicon Valley Chapter Organizer, *The AI Collective* Jun 2025 – Present

- Designed and directed a five-event RAG and AI agents enablement program with SupportVectors AI Labs — setting the topics, direction, and course structure — reaching 400+ attendees and increasing enrollment 38%.
- Organized and moderated an [Enterprise RAG fireside at Snowflake Menlo Park](#) for 80+ senior technical leaders, sourcing and hosting a panel of a CTO, Senior Directors, and a Principal Engineer from Snowflake, Zscaler, and SupportVectors.

CORE TECHNICAL SKILLS

Applied AI & Agents: LLM Applications, Retrieval-Augmented Generation (RAG), Agentic Workflows, Multi-Agent Systems, Hybrid Retrieval, Multimodal AI, Prompt Engineering, Context Engineering, MCP, Fine-Tuning

Reliability & Validation: Grounding, Hallucination Mitigation, Source Validation, Retrieval Precision Measurement, Reranking, Statistical Validation, Uncertainty Quantification

AI Frameworks: LangGraph, LangChain, DSPy, CrewAI, Hugging Face, PyTorch, TensorFlow, scikit-learn

LLM Platforms: Claude, OpenAI (GPT-4o), Ollama, Mistral

Retrieval & Data: Elasticsearch, Qdrant, BM25, Dense Embeddings, Reciprocal Rank Fusion, PostgreSQL, MySQL, Pandas, NumPy

AI Application Engineering: Python, FastAPI, REST APIs, OAuth, Next.js, React, Streamlit

Cloud & Production: Docker, Kubernetes, Git/GitHub, CI/CD, AWS, GCP, Azure

SELECTED APPLIED AI PROJECTS

Agentic Research & Content Intelligence System (CrewAI, NLP, Multi-Agent Systems):

- Developed an early multi-agent framework — coordinating 6+ specialized agents across search, validation, analysis, and content generation — to test reliability patterns ahead of mainstream agentic-AI adoption.
- Separated research and validation stages — using diverse search strategies and source-validation checks — to mitigate hallucinations and reduce unsupported or misleading content.
- Implemented agent behavior using a Belief-Desire-Intention (BDI) architecture to separate information gathering, reasoning, validation, and content-generation responsibilities.

Privacy-Aware Multimodal Search & Annotation (Qdrant, OpenCV, DeepFace, GPT-4o, Streamlit):

- Built a multimodal AI workflow to identify, retrieve, and anonymize people across image collections, cutting manual annotation effort 70%.
- Combined RetinaFace detection and ArcFace embeddings with vector similarity search, achieving 98% retrieval consistency; designed a Qdrant metadata architecture with optimized batch operations that increased search throughput $\sim 3\times$.

TECHNICAL WRITING & PUBLICATIONS

[DRAGIN](#): Dynamic Retrieval Augmented Generation Based on the Information Needs of LLMs

[From Assumptions to Accuracy](#): The Role of Conditional Probability in Real-World Predictions

CERTIFICATIONS

Google Cloud Certified: Professional Machine Learning Engineer

SupportVectors AI Labs: RAG and AI Search, AI Agents Bootcamp, LLM Bootcamp, Foundations of AI and Neural Networks

Society of Actuaries (SOA): Probability, Financial Mathematics, Financial Engineering, Statistical Risk Modeling, Mathematical Statistics

EDUCATION

Georgetown University

Master of Science, Data Science and Analytics

Washington, D.C.

Mount Carmel College

Bachelor of Science, Statistics and Mathematics

Bengaluru, India